

```
In [50]: import pandas as pd
```

```
In [51]: df = pd.read_csv("Social_Network_Ads.csv")
```

```
In [52]: df
```

```
Out[52]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

```
In [53]: x=df[['Age','EstimatedSalary']]
```

```
In [54]: y=df['Purchased']
```

```
In [55]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
x_scaled = scaler.fit_transform(x)
```

```
In [56]: from sklearn.model_selection import train_test_split
```

```
In [117... x_train,x_test,y_train,y_test=train_test_split(x_scaled,y,random_state=0, te
```

```
In [118... x_train
```

```
Out[118... array([[0.61904762, 0.17777778],
 [0.33333333, 0.77777778],
 [0.47619048, 0.25925926],
 [0.33333333, 0.88888889],
 [0.80952381, 0.04444444],
 [0.83333333, 0.65925926],
 [0.5, 0.2],
 [0.47619048, 0.34074074],
 [0.42857143, 0.25925926],
 [0.42857143, 0.35555556],
 [0.4047619, 0.07407407],
 [0.4047619, 0.25925926],
 [0.57142857, 0.42962963],
 [0.69047619, 0.25185185],
 [0.97619048, 0.1037037],
 [0.73809524, 0.37037037],
 [0.64285714, 0.85925926],
 [0.30952381, 0.54814815],
 [0.66666667, 0.4962963],
 [0.69047619, 0.26666667],
 [0.19047619, 0.],
 [1., 0.64444444],
 [0.47619048, 0.71851852],
 [0.52380952, 0.68148148],
 [0.57142857, 0.28148148],
 [0.4047619, 0.32592593],
 [0.71428571, 0.19259259],
 [0.71428571, 0.88148148],
 [0.47619048, 0.72592593],
 [0.26190476, 0.98518519],
 [0.19047619, 0.],
 [1., 0.2],
 [0.14285714, 0.02962963],
 [0.57142857, 0.99259259],
 [0.66666667, 0.6],
 [0.23809524, 0.32592593],
 [0.5, 0.6],
 [0.23809524, 0.54814815],
 [0.54761905, 0.42222222],
 [0.64285714, 0.08148148],
 [0.35714286, 0.4],
 [0.04761905, 0.4962963],
 [0.30952381, 0.43703704],
 [0.57142857, 0.48148148],
 [0.4047619, 0.42222222],
 [0.35714286, 0.99259259],
 [0.52380952, 0.41481481],
 [0.78571429, 0.97037037],
 [0.66666667, 0.47407407],
 [0.4047619, 0.44444444],
 [0.47619048, 0.26666667],
 [0.42857143, 0.44444444],
 [0.45238095, 0.46666667],
 [0.47619048, 0.34074074],
 [1., 0.68888889],
 [0.04761905, 0.4962963],
```

```
[0.73809524, 0.17777778],  
[0.21428571, 0.11851852],  
[0.02380952, 0.40740741],  
[0.5, 0.47407407],  
[0.19047619, 0.48888889],  
[0.16666667, 0.48148148],  
[0.23809524, 0.51851852],  
[0.88095238, 0.17777778],  
[0.76190476, 0.54074074],  
[0.73809524, 0.54074074],  
[0.80952381, 1.],  
[0.4047619, 0.37037037],  
[0.57142857, 0.28888889],  
[0.38095238, 0.20740741],  
[0.45238095, 0.27407407],  
[0.71428571, 0.11111111],  
[0.26190476, 0.20740741],  
[0.42857143, 0.27407407],  
[0.21428571, 0.28888889],  
[0.19047619, 0.76296296]])
```


In [119... `y_train`

```
Out[119... 250    0  
63      1  
312     0  
159     1  
283     1  
..  
323     1  
192     0  
117     0  
47      0  
172     0  
Name: Purchased, Length: 300, dtype: int64
```

In [120... `from sklearn.linear_model import LogisticRegression`

In [121... `import seaborn as sns
sns.countplot(x=y)`

Out[121... `<Axes: xlabel='Purchased', ylabel='count'>`



```
In [122... y.value_counts()
```

```
Out[122... Purchased
0      257
1      143
Name: count, dtype: int64
```

```
In [123... classifier = LogisticRegression()
```

```
In [124... classifier.fit(x_train,y_train)
```

```
Out[124... ▼ LogisticRegression ⓘ ?
LogisticRegression()
```

```
In [125... y_pred = classifier.predict(x_test)
```

```
In [126... y_train.shape
```

```
Out[126... (300,)
```

```
In [127... x_train.shape
```

```
Out[127... (300, 2)
```

```
In [128... y_pred
```

```
Out[128...] array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
        0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
        0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1,
        0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1])
```


```
In [129...] y_test
```

```
Out[129...] 132    0
            309    0
            341    0
            196    0
            246    0
            ..
            146    1
            135    0
            390    1
            264    1
            364    1
            Name: Purchased, Length: 100, dtype: int64
```

```
In [130...] import matplotlib.pyplot as plt
```

```
In [131...] plt.xlabel('Age')
            plt.ylabel('Salary')
            plt.scatter(x['Age'],x['EstimatedSalary'],c=y)
```

```
Out[131...] <matplotlib.collections.PathCollection at 0x7fe9d4927fd0>
```




```
In [132... from sklearn.preprocessing import MinMaxScaler  
scaler = MinMaxScaler()  
x_scaled = scaler.fit_transform(x)
```

```
In [133... pd.DataFrame(x_scaled).describe()
```

```
Out[133...      0      1  
count  400.000000  400.000000  
mean    0.467976  0.405500  
std     0.249592  0.252570  
min     0.000000  0.000000  
25%    0.279762  0.207407  
50%    0.452381  0.407407  
75%    0.666667  0.540741  
max     1.000000  1.000000
```

```
In [134... plt.xlabel('Age')  
plt.ylabel('Salary')  
plt.scatter(x_scaled[:,0],x_scaled[:,1],c=y)
```

```
Out[134... <matplotlib.collections.PathCollection at 0x7fe9d4999e50>
```



```
In [135... from sklearn.metrics import confusion_matrix
```

```
In [136... confusion_matrix(y_test,y_pred)
```

```
Out[136... array([[67,  1],
          [10, 22]])
```

```
In [137... y_pred
```

```
Out[137... array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
        0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
        0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1,
        0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1])
```


```
In [138... y_test.value_counts()
```

```
Out[138... Purchased
0      68
1      32
Name: count, dtype: int64
```

```
In [139... from sklearn.metrics import ConfusionMatrixDisplay
```

```
In [140... ConfusionMatrixDisplay.from_estimator(classifier, x_test, y_test)
```

```
Out[140... <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7fe9d4b0aa50>
```



```
In [141... from sklearn.metrics import accuracy_score
```

```
accuracy_score(y_test,y_pred)
```

Out[141...] 0.89

```
In [142...] from sklearn.metrics import classification_report
```

```
In [143...] print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.99	0.92	68
1	0.96	0.69	0.80	32
accuracy			0.89	100
macro avg	0.91	0.84	0.86	100
weighted avg	0.90	0.89	0.88	100